

Librerias

```

1 from google.colab import drive
2 drive.mount('/content/drive')

    Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

1 import numpy as np
2 import matplotlib.pyplot as plt
3 import tensorflow as tf
4 import os
5 from keras.src.engine.data_adapter import train_validation_split

```

Procesamiento de los datos

```

1 BATCH_SIZE = 32
2 IMG_SIZE = (160,160)
3 PATH = '/content/drive/MyDrive/Images/computer_parts'
4
5 train_dir = os.path.join(PATH, 'train')
6
7 train_dataset = tf.keras.utils.image_dataset_from_directory(train_dir,
8                                                             shuffle = True,
9                                                             batch_size = BATCH_SIZE,
10                                                            image_size = IMG_SIZE)
11
12 test_dir = os.path.join(PATH, 'test')
13
14 test_dataset = tf.keras.utils.image_dataset_from_directory(train_dir,
15                                                            shuffle = True,
16                                                            batch_size = BATCH_SIZE,
17                                                            image_size = IMG_SIZE)

    Found 853 files belonging to 3 classes.
    Found 853 files belonging to 3 classes.

1 train_dataset.take(1)

<_TakeDataset element_spec=(TensorSpec(shape=(None, 160, 160, 3), dtype=tf.float32, name=None), TensorSpec(shape=(None,),
dtype=tf.int32, name=None))>

1 class_name = train_dataset.class_names
2
3 plt.figure(figsize = (10,10))
4 for image, label in train_dataset.take(1):
5     for i in range(9):
6         ax=plt.subplot(3,3,i+1)
7         plt.imshow(image[i].numpy().astype("uint8"))
8         plt.title(class_name[label[i]])
9         plt.axis("off")

```

monitor



keyboard



keyboard



keyboard



mouse



keyboard



```
1 val_batches = tf.data.experimental.cardinality(train_dataset)
2 test_batches = tf.data.experimental.cardinality(test_dataset)
3 validation_dataset = train_dataset.skip(val_batches // 5)
4
5 print('Numero de batches para validation_dataset = %d' % tf.data.experimental.cardinality(validation_dataset))
6 print('Numero de batches para test_dataset = %d' % test_batches)
```

```
Numero de batches para validation_dataset = 22
Numero de batches para test_dataset = 27
```






```
1 AUTOTUNE = tf.data.AUTOTUNE
2
3 test_dataset = test_dataset.prefetch(buffer_size=AUTOTUNE)
4 train_dataset = train_dataset.prefetch(buffer_size=AUTOTUNE)
5 validation_dataset = validation_dataset.prefetch(buffer_size=AUTOTUNE)
```



```
1 data_augmentation = tf.keras.Sequential([
2     tf.keras.layers.RandomFlip('horizontal'),
3     tf.keras.layers.RandomRotation(0.2)
4 ])
```

```
1 for image, _ in train_dataset.take(1):
2     plt.figure(figsize=(10,10))
3     first_image = image[0]
4     for i in range(9):
5         ax=plt.subplot(3,3,i+1)
6         augmented_image = data_augmentation(tf.expand_dims(first_image,0))
7         plt.imshow(augmented_image[0] / 255)
8         plt.axis("off")
```



```
1 rescale = tf.keras.layers.Rescaling(1./127.5, offset=-1)
2 preprocess_input = tf.keras.applications.mobilenet_v2.preprocess_input
```

Cargar la red pre-entrenada DenseNet121



```
1 IMG_SHAPE = IMG_SIZE + (3,)
2 print(IMG_SHAPE)
3 base_model = tf.keras.applications.DenseNet121(input_shape=IMG_SHAPE,
4                                                include_top=False,
5                                                weights='imagenet')
```

```
(160, 160, 3)
```



```
1 image_batch, label_batch = next(iter(train_dataset))
2 feature_batch = base_model(image_batch)
3
4 print(feature_batch.shape)
```

```
(32, 5, 5, 1024)
```

```
1 base_model.trainable = False
2 base_model.summary()
```

atenate)			conv5_block15_2_conv[0][0]
conv5_block16_0_bn (Batch Normalization)	(None, 5, 5, 992)	3968	['conv5_block15_concat[0][0]']
conv5_block16_0_relu (Activation)	(None, 5, 5, 992)	0	['conv5_block16_0_bn[0][0]']
conv5_block16_1_conv (Conv 2D)	(None, 5, 5, 128)	126976	['conv5_block16_0_relu[0][0]']
conv5_block16_1_bn (Batch Normalization)	(None, 5, 5, 128)	512	['conv5_block16_1_conv[0][0]']
conv5_block16_1_relu (Activation)	(None, 5, 5, 128)	0	['conv5_block16_1_bn[0][0]']
conv5_block16_2_conv (Conv 2D)	(None, 5, 5, 32)	36864	['conv5_block16_1_relu[0][0]']
conv5_block16_concat (Concatenate)	(None, 5, 5, 1024)	0	['conv5_block15_concat[0][0]', 'conv5_block16_2_conv[0][0]']
bn (Batch Normalization)	(None, 5, 5, 1024)	4096	['conv5_block16_concat[0][0]']
relu (Activation)	(None, 5, 5, 1024)	0	['bn[0][0]']
=====			
Total params: 7037504 (26.85 MB)			
Trainable params: 0 (0.00 Byte)			
Non-trainable params: 7037504 (26.85 MB)			

```
1 global_average_layer = tf.keras.layers.GlobalAveragePooling2D()
2 feature_batch_average = global_average_layer(feature_batch)
3 print(feature_batch_average.shape)
```

(32, 1024)

```
1 prediction_layer = tf.keras.layers.Dense(3)
2 prediction_batch = prediction_layer(feature_batch_average)
3 print(prediction_batch.shape)
```

(32, 3)

Unir modelo

```
1 inputs = tf.keras.Input(shape=(160,160,3))
2 x = data_augmentation(inputs)
3 x = preprocess_input(x)
4 x = base_model(x, training=False)
5 x = global_average_layer(x)
6 x = tf.keras.layers.Dropout(0.2)(x)
7 outputs = prediction_layer(x)
8 model = tf.keras.Model(inputs, outputs)
9

1 base_learning_rate = 0.0001
2 model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=base_learning_rate),
3               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
4               metrics=['accuracy'])
5 model.summary()
```

Model: "model_7"		
Layer (type)	Output Shape	Param #
=====		
input_16 (InputLayer)	[(None, 160, 160, 3)]	0
sequential_7 (Sequential)	(None, 160, 160, 3)	0
tf.math.truediv_7 (TFOPLambda)	(None, 160, 160, 3)	0
tf.math.subtract_7 (TFOPLambda)	(None, 160, 160, 3)	0

densenet121 (Functional)	(None, 5, 5, 1024)	7037504
global_average_pooling2d_7 (GlobalAveragePooling2D)	(None, 1024)	0
dropout_7 (Dropout)	(None, 1024)	0
dense_7 (Dense)	(None, 3)	3075

```

=====
Total params: 7040579 (26.86 MB)
Trainable params: 3075 (12.01 KB)
Non-trainable params: 7037504 (26.85 MB)

```

```

1 initial_epochs = 10
2
3 loss0, accuracy0 = model.evaluate(validation_dataset)
4
5 print(loss0)
6 print(accuracy0)

22/22 [=====] - 11s 187ms/step - loss: 0.9614 - accuracy: 0.6537
0.9613553881645203
0.6536796689033508

```

```

1 history = model.fit(train_dataset,
2                       epochs=initial_epochs,
3                       validation_data = validation_dataset)

Epoch 1/10
27/27 [=====] - 17s 291ms/step - loss: 0.9987 - accuracy: 0.6237 - val_loss: 0.6455 - val_accuracy: 0.7359
Epoch 2/10
27/27 [=====] - 10s 318ms/step - loss: 0.8429 - accuracy: 0.6542 - val_loss: 0.5013 - val_accuracy: 0.8124
Epoch 3/10
27/27 [=====] - 10s 331ms/step - loss: 0.6840 - accuracy: 0.7140 - val_loss: 0.4281 - val_accuracy: 0.8384
Epoch 4/10
27/27 [=====] - 9s 296ms/step - loss: 0.5972 - accuracy: 0.7655 - val_loss: 0.3529 - val_accuracy: 0.8846
Epoch 5/10
27/27 [=====] - 10s 302ms/step - loss: 0.4846 - accuracy: 0.8159 - val_loss: 0.3126 - val_accuracy: 0.9091
Epoch 6/10
27/27 [=====] - 10s 332ms/step - loss: 0.4618 - accuracy: 0.8242 - val_loss: 0.2804 - val_accuracy: 0.9235
Epoch 7/10
27/27 [=====] - 10s 320ms/step - loss: 0.4391 - accuracy: 0.8406 - val_loss: 0.2438 - val_accuracy: 0.9394
Epoch 8/10
27/27 [=====] - 9s 301ms/step - loss: 0.4014 - accuracy: 0.8605 - val_loss: 0.2164 - val_accuracy: 0.9466
Epoch 9/10
27/27 [=====] - 11s 344ms/step - loss: 0.3416 - accuracy: 0.8746 - val_loss: 0.1954 - val_accuracy: 0.9567
Epoch 10/10
27/27 [=====] - 10s 348ms/step - loss: 0.3214 - accuracy: 0.8863 - val_loss: 0.1787 - val_accuracy: 0.9553


```

```

1 acc = history.history['accuracy']
2 val_acc = history.history['val_accuracy']
3 plt.figure(figsize=(8,8))
4 plt.plot(acc, label = 'Training acc')
5 plt.plot(val_acc, label = 'Validation acc')

```

[<matplotlib.lines.Line2D at 0x7df24e152e60>]



Fine tuning

```
1 base_model.trainable = True
2
3 print('Numero de capas ', len(base_model.layers))

Numero de capas 427

1 fine_tune_at=100
2 for layer in base_model.layers[:fine_tune_at]:
3     layer.trainable = False

1 model.compile(optimizer=tf.keras.optimizers.RMSprop(learning_rate=base_learning_rate),
2               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
3               metrics=['accuracy'])
4 model.summary()
```

Model: "model_7"

Layer (type)	Output Shape	Param #
=====		
input_16 (InputLayer)	[(None, 160, 160, 3)]	0
sequential_7 (Sequential)	(None, 160, 160, 3)	0
tf.math.truediv_7 (TFOpLam bda)	(None, 160, 160, 3)	0
tf.math.subtract_7 (TFOpLam mbda)	(None, 160, 160, 3)	0
densenet121 (Functional)	(None, 5, 5, 1024)	7037504
global_average_pooling2d_7 (GlobalAveragePooling2D)	(None, 1024)	0
dropout_7 (Dropout)	(None, 1024)	0
dense_7 (Dense)	(None, 3)	3075
=====		
Total params: 7040579 (26.86 MB)		
Trainable params: 6152259 (23.47 MB)		
Non-trainable params: 888320 (3.39 MB)		

```
1 fine_tune_epochs = 10
2 total_epochs = initial_epochs + fine_tune_epochs
3 history_fine = model.fit (train_dataset,
4                           epochs = total_epochs,
5                           initial_epoch = history.epoch[-1],
6                           validation_data=validation_dataset)
```

Epoch 10/20
27/27 [=====] - 55s 511ms/step - loss: 0.5351 - accuracy: 0.8581 - val_loss: 0.0598 - val_accuracy: 0.9827

```


Epoch 11/20
27/27 [=====] - 11s 374ms/step - loss: 0.1131 - accuracy: 0.9613 - val_loss: 0.0222 - val_accuracy: 0.9942
Epoch 12/20
27/27 [=====] - 10s 347ms/step - loss: 0.0662 - accuracy: 0.9730 - val_loss: 0.0385 - val_accuracy: 0.9856
Epoch 13/20
27/27 [=====] - 11s 367ms/step - loss: 0.0916 - accuracy: 0.9707 - val_loss: 0.0054 - val_accuracy: 0.9986
Epoch 14/20
27/27 [=====] - 11s 375ms/step - loss: 0.1060 - accuracy: 0.9730 - val_loss: 0.0230 - val_accuracy: 0.9928
Epoch 15/20
27/27 [=====] - 11s 348ms/step - loss: 0.0595 - accuracy: 0.9812 - val_loss: 0.1460 - val_accuracy: 0.9437
Epoch 16/20
27/27 [=====] - 11s 361ms/step - loss: 0.0889 - accuracy: 0.9801 - val_loss: 0.0089 - val_accuracy: 0.9986
Epoch 17/20
27/27 [=====] - 11s 368ms/step - loss: 0.0225 - accuracy: 0.9930 - val_loss: 0.0045 - val_accuracy: 0.9986
Epoch 18/20
27/27 [=====] - 11s 359ms/step - loss: 0.0264 - accuracy: 0.9894 - val_loss: 9.5069e-04 - val_accuracy: 1.
Epoch 19/20
27/27 [=====] - 10s 324ms/step - loss: 0.0315 - accuracy: 0.9906 - val_loss: 0.0295 - val_accuracy: 0.9913
Epoch 20/20
27/27 [=====] - 11s 374ms/step - loss: 0.0111 - accuracy: 0.9965 - val_loss: 0.0045 - val_accuracy: 0.9986

```

```

1 acc += history_fine.history['accuracy']
2 val_acc += history_fine.history['val_accuracy']
3
4 plt.figure(figsize=(8,8))
5 plt.plot(acc, label = 'training acc')
6 plt.plot(val_acc, label = 'validation acc')
7 plt.ylim([0.5, 1])
8 plt.plot([initial_epochs, initial_epochs],
9          plt.ylim(), label='start fine tuning')
10
11 plt.legend(loc='lower right')
12 plt.xlabel('epochs')
13 plt.show()

```



```
1 model.save('model.h5')
```

```

1 predictions = model.predict(test_dataset)
2
3 # class_name = train_dataset.class_names
4
5 plt.figure(figsize = (10,10))
6 for image, label in train_dataset.take(1):
7     for i in range(9):
8         ax=plt.subplot(3,3,i+1)
9         plt.imshow(image[i].numpy().astype("uint8"))
10        plt.title(class_name[np.argmax(predictions[i])])
11        plt.axis("off")

```

27/27 [=====] - 6s 176ms/step

keyboard



keyboard



monitor



mouse



mouse



monitor



monitor



mouse



mouse



